

How does the latest Prophet forecasting model compare to the ARIMA model? Evaluation of BTC/ZAR predictability through a Mincer-Zarnowitz approach

14th September 2017

Fiona Ganie (GNXFIO001)

GNXFIO001@myuct.ac.za

Abstract

The Autoregressive Integrated Moving Average (ARIMA) model is one of the most widely used time series models that has gained popularity in the exchange rate market due to its ease of implementation and tractability. With the evolution of computational power, soft computing techniques have since been used in exchange rate markets. The ARIMA model has been used as the standard benchmark model against which these complex methods have been compared. Although they have produced significantly better results than the ARIMA model, these models lack interpretability and building them is a challenging task. Prophet is a sophisticated model that differs to traditional time series models in that it can be utilised by non-experts who have little knowledge about the statistical intricacies involved in the model, however have domain knowledge about the data generating process. Prophet is also robust to outliers and can handle missing values in the time series without the need for interpolation. This paper compares the out-of-sample forecasts of Prophet and ARIMA over varying forecast horizons, using a Mincer-Zarnowitz approach for

forecast evaluation. A complete dataset of the prices of Bitcoin/Rand as well as a dataset which has missing prices and outliers present is used in the analysis. It is found that the ARIMA model produces the most efficient and unbiased forecasts over all forecast horizons, and that the ARIMA model is more robust when faced with missing values and outliers present in the data. Keywords: Prophet, ARIMA, Mincer-Zarnowitz, Box-Jenkins, Bitcoin.

1 Introduction

The Autoregressive Integrated Moving Average (ARIMA) model is one of the most widely used time series models that has attracted attention in financial market forecasting (Khashei, Bijari, and Ardali 2009). Although the Random Walk model has typically been applied to foreign exchange markets and has produced superior results, some researchers have contended that foreign exchange markets are not efficient and believe that future prices depend on current and past events (Abu-Mostafa and Atiya 1996). This has led to the application of the ARIMA model to exchange rate problems, where it has since gained popularity due to its ease of implementation and tractability. The ARIMA model is easy to use and yields good forecasts over short forecast horizons, however the linearity of the ARIMA model fails to adequately capture the non-linearity inherent in exchange rate data (Zhang and Hu 1998).

With the evolution of computational power, non-linear, soft computing techniques have been proposed as a solution. The ARIMA model has been used as the standard benchmark model against which these more complex methods have been compared. Although they have produced significantly better results than the ARIMA model, these models lack interpretability and building them is a challenging task. The ARIMA model is tractable and less computationally expensive. It has been used as the building blocks for more advanced models and has provided the inspiration for hybrid versions of the model which have been used in exchange rate forecasting.

In December 2016, Facebook open-sourced their forecasting model Prophet. Prophet differs to traditional time series model such as the ARIMA, in that it can produce high quality forecasts in a straightforward way. Prophet is a sophisticated model that provides informative results, however is config-

urable and easy to use. It can be utilised by non-experts who have little knowledge about the statistical intricacies involved in the model, however have domain knowledge about the data generating process. This knowledge can be easily incorporated into the model through its intuitively adjustable parameters (Taylor and Letham 2017). Furthermore, Prophet is robust to outliers and can handle missing values in the time series without the need for interpolation (Taylor and Letham 2017). This allows analysts without knowledge on how to pre-process the data, to utilise the model without sacrificing predictive accuracy. If Prophet produces forecasts that are as good as the ARIMA model, it can be compared to more complex forecasting methods that are less tractable and flexible.

This paper broadly aims to compare the forecasts produced by the traditional ARIMA model and Prophet through an evaluation of the exchange rate between Bitcoin and the Rand (BTC/ZAR), using a Mincer-Zarnowitz approach of measuring forecast accuracy. More specifically, this paper aims to:

1. Compare the out-of-sample forecasts of ARIMA and Prophet over varying forecast horizons.
2. Compare the out-of-sample forecast performance with missing values and outliers present in the data.

Section 2 will present the key findings from the major works in which comparisons have been made to the ARIMA model in exchange rate forecasting, as well as the theory behind the ARIMA and Prophet model. Section 3 will provide a brief explanation and motivation for the dataset used in this paper. Section 4 will describe the methodology used in selecting the ARIMA and Prophet models and the Mincer-Zarnowitz approach of measuring fore-

cast accuracy. The results from applying the ARIMA and Prophet model to the dataset is presented in Section 5. Finally, Section \ref{Discussion and Conclusions} discusses the results in a statistical context.

2 Background and Theory

2.1 Forecasting Exchange Rates with the ARIMA model

ARIMA models became highly popular since its introduction by Box and Jenkins, when it was shown that they could outperform complex econometric models in a variety of situations (Hibon and Makridakis 1997). The ARIMA model expresses the process $\{y_t\}$ as a function of the weighted average of past values of the process and lagged values of the residuals. The weighted average of the past p values of the process represents an autoregressive (AR) process of order p . It feeds back past values of the process into the current value, inducing correlation between all lags of the process. The weighted average of the q lagged residuals represents a moving average (MA) process of order q . The purpose of mixing the MA process with the AR process is to reduce the large number of past values required by AR processes and to control for the autocorrelation which it creates between lagged values of the process. The combination of the AR(p) and MA(q) process results in a more parsimonious model, and forms a stationary autoregressive moving average (ARMA(p,q)) process defined as:

$$y_t = c + \sum_{i=1}^p \theta_i y_{t-i} + \sum_{i=1}^q \phi_i e_{t-i} + e_t$$

where c is a constant and $\{e_t\}$ is a white noise process with zero mean and variance σ^2 .

The ARIMA(p,d,q) model generalises the ARMA model in that it includes both stationary and non-stationary processes. The parameter d is the degree of differencing required to render the process stationary. If d is equal to zero the process is stationary and equivalent to an ARMA model, and if d is strictly positive the process requires differencing to make it stationary. The ARIMA model can be defined succinctly using the backward shift operator B, which shifts the process back by one unit of time, and is defined as $By_t = y_{t-1}$. The ARIMA model has the form:

$$(1 - \sum_{i=1}^p \theta_i B^i)(1 - B)^d y_t = c + (1 + \sum_{i=1}^q \phi_i B^i) e_t$$

where c is a constant and $\{e_t\}$ is a white noise process with zero mean and variance σ^2 .

ARIMA models have commonly been used in financial forecasting and are popular for observing stock prices and exchange rates due to its power and statistical properties (C.-S. Lin, Chiu, and Lin 2012). They have frequently been used as a benchmark to compare new forecasting techniques that have emerged over time. Kamruzzaman and Sarker (2003) applied an ARIMA model to forecast the exchange rate between the Australian dollar and six other currencies, and used the forecast errors as well as the accuracy of the direction of the forecasts to evaluate the performance of three Neural Network models. In a similar study, Khashei, Bijari, and Ardali (2012) compared the predictive capability of their proposed hybrid model consisting of an ARIMA and a Probabilistic Neural Network (PNN) to the traditional ARIMA model, and justified the use of the ARIMA as a benchmark by claiming that it is the most important linear model. Their comparison was made by investigating the forecasts between the British pound and US dollar over varying forecast horizons. Not only has the ARIMA model been useful as a benchmark, but

it has also yielded satisfactory results when predicting exchange rates. When Nwankwo (2014) forecasted the rate between the Nairo and dollar, diagnostic testing revealed that the ARIMA(1,0,0) model was the best fit for the data based on Akaike's Information Criterion (AIC).

The ARIMA model is attractive as it is tractable and produces good short-term forecasts when more than 100 observations are used (Tseng et al. 2001). Although the ARIMA model has the advantage of ease of implementation and flexibility, it fails to capture the non-linearity and volatility present in exchange rate data. Over time, the ARIMA model has evolved to cater for a wider variety of data and to compensate for some of its shortcomings. The most popular versions of the ARIMA model that has been implemented in exchange rate forecasting is the Seasonal ARIMA (SARIMA) and Fractional ARIMA model (FARIMA).

The SARIMA model was introduced to capture the periodic behaviour of data and extends the ARIMA class by including a term for seasonal differencing. Etuk, Wokoma, and Moffat (2013) modelled the Naira/CFA Franc exchange rate which exhibited monthly seasonality using an additive SARIMA model, to demonstrate that it can be a useful fit for exchange rate data which displays seasonality. Their results showed that the SARIMA model adequately described the variation in the exchange rate series. The ARFIMA model generalises the ARIMA model in that the degree required to make the data stationary can assume any real value, and is no longer restricted to the integer domain. The ARFIMA model has the ability to capture the dependence between observations that are widely spread apart in time (Cheung 1993). This makes the model more parsimonious since it can capture long memory in data as well as short term dynamics (Cheung 1993). Cheung (1993) fitted the ARFIMA model to examine five exchange rates, and found

that there was strong evidence of long memory in the exchange rate time series.

Generalised autoregressive conditional heteroskedasticity (GARCH) models were later developed in an attempt to capture the volatility in financial markets (Anastasakis and Mort 2009). Hsieh (1989) applied a GARCH model to investigate five exchange rates and his results showed that although the GARCH model outperformed the random walk, some non-linear information still remained in the residuals.

2.2 Other techniques used to forecast exchange rates

Over time, financial forecasting methods have moved away from linear models like ARIMA and GARCH, to soft computing techniques. These more complex techniques are non-linear and can fit complex time series more easily (Castillo and Melin 2002). Unlike the ARIMA model, soft computing techniques do not impose structural assumptions on the model apriori (Castillo and Melin 2002). Some of the most commonly used artificial intelligence methods used to forecast exchange rate data are Neural Networks and Fuzzy Logistic Systems.

Artificial Neural Networks (ANNs) have had many successful applications in forecasting exchange rates and are more advantageous than other non-linear forecasting methods. They are data driven, can adapt to non-stationary environments and can approximate any continuous function (Khashei and Bijari 2011). In a study done by Fahimifard et al. (2009), the ANN was found as an effective way to improve the forecasts of exchange rates. Superior results were produced when compared to the ARIMA and GARCH model using the root mean square error (RMSE), mean square error (MSE) and mean absolute difference as a measure of performance. Fuzzy Logistic Sys-

tems (FLSs) were initially developed to solve problems involving linguistic terms, and have been successfully used in financial forecasting (Khashei, Bijari, and Ardali 2009). Fuzzy logic tries to imitate human reasoning and the decision-making process and allows for finer rather than discrete decisions to be provided. Santos, Costa, and Santos Coelho (2007) investigated how well FLSs and ANNs perform compared to the traditional ARMA and GARCH model. He examined the forecasts of Brazilian exchange rate returns by considering different frequencies of the series and comparing their one step-ahead forecasts. By analysing accuracy statistics, he found that FLSs and ANNs achieved higher returns based on the forecasts they produced. Similar results were found by Khashei, Bijari, and Ardali (2009) when he analysed the predictive capabilities of FLSs, ANNs, the traditional ARIMA model, and a Fuzzy ARIMA model.

Although ANNs have been broadly applied in financial forecasting, the process of building them is a complex task and there is no consistent method of design compared to the traditional Box-Jenkins ARIMA model. Unlike ARIMA models, the performance of ANNs is sensitive to many modelling factors such as the number of input nodes included and the size of the training sample chosen (Zhang and Hu 1998). Like ANNs there is no systematic approach to designing FLSs and they are only understandable when simple. Although FLSs has the advantage over ARIMA models that they can be applied to data with few observations available, it gives acceptable rather than accurate results and are more suitable for problems which do not require high accuracy.

2.3 Hybrid ARIMA Models

Over time, many researchers began to think of ways in which to harness the advantages of tractable linear models such as the ARIMA model, and of more complex non-linear models. By combining the ARIMA model with other forecasting methods, the advantages of both forecasting methods are leveraged while simultaneously improving their limitations. Some of the hybrid models which have commonly been used in exchange rate forecasting are the Fuzzy ARIMA and ANN-ARIMA model.

The ARIMA model produces very accurate forecasts over short time horizons however it has the limitation of requiring more than 100 observations of historical data to yield accurate results (Tseng et al. 2001). In a world that is constantly changing and with the rapid advancement of technology, access to large amounts of historical data is difficult to obtain. On the contrary, a fuzzy regression model requires little historical data however produces wide prediction intervals if extreme values are present in the data. The Fuzzy ARIMA model combines the ARIMA and fuzzy regression model to exploit the advantages of both models while simultaneously overcoming their limitations. Tseng et al. (2001) proposed applying a Fuzzy ARIMA model to forecast the exchange rate of Taiwan dollars to US dollars to demonstrate the model's appropriateness and power. The Fuzzy ARIMA not only produced forecasts that were superior to the ARIMA and fuzzy time series models but also provided an upper and lower bound which can be used by decision makers to determine the best and worst possible situations.

Although ARIMA models are powerful, they require non-stationary data to be differenced and impose prior assumptions onto the distribution of the data (Ince and Trafalis 2006). In contrast, machine learning techniques such as ANNs do not impose any assumptions onto the data generating process

however, being data-driven, are sensitive to the number of input nodes used. The ANN-ARIMA model draws on the strengths of both models and overcomes these individual difficulties. Ince and Trafalis (2006) created an ANN-ARIMA model by using the ARIMA model to determine the number input nodes required by an ANN for three exchange rates. When the ANN-ARIMA model was compared to the pure ARIMA model, the hybrid model outperformed the ARIMA based on the MSE.

Although the traditional ARIMA models produce less superior forecasts than its hybrid forms and other complex non-linear techniques, its forecasts are still satisfactory. They are simple models that are easy to implement and have a consistent method of model design and selection. ARIMA models are also more robust and efficient than complex structural models in relation to short-run forecasting. The fact that they have been used as the foundation for more advanced models and have commonly been used as a benchmark for comparison, justifies it as a good starting point to compare it to Facebook’s forecasting method, Prophet, that was recently released.

2.4 Forecasting with Prophet

The techniques that have been considered for exchange rate forecasting thus far, require the analyst to have vocational knowledge about time series. Prophet differs to traditional time series models in that it is flexible and can be customised by a large number of non-experts who have little knowledge about time series, however have domain knowledge about the data generating process. Prophet allows for a large number of forecasts to be produced across a variety of problems and consists of a robust evaluation system that allows for a large number of forecasts be evaluated and compared. This is Facebook’s definition of forecasting at scale. Prophet is similar to a Gener-

alized Additive Model (GAM) - an additive regression model that consists of non-linear and linear regression functions applied to predictor variables (Taylor and Letham 2017). The decomposable model is of the form:

$$y(t) = g(t) + s(t) + h(t) + e_t$$

where the components of the model represent the growth, seasonality and holiday respectively, and e_t is white noise.

Prophet, like the GAM, frames the forecasting problem as a curve fitting exercise and uses backfitting to find the regression functions. This allows for the model to be fitted quickly and missing values and large outliers to be handled elegantly. The regression model also provides model flexibility and allows the analyst to interactively change model parameters (Taylor and Letham 2017).

The growth component of Prophet may be modelled as a linear or non-linear function of time. Linear growth is modelled by a piecewise constant function while non-linear growth is modelled similar to population growths which use a logistic growth model (Taylor and Letham 2017). A time-varying upper limit may be specified for logistic growth, at which point the forecasts will saturate. This carrying capacity allows the analyst to incorporate their prior knowledge about the maximum obtainable growth level such as the total market or population size into the model. Prophet accounts for changes in the trajectory of this trend by automatically detecting and selecting change-points in the data at which the growth rate is allowed to change. These changepoints have a Laplace prior distribution placed on them and its scale parameter may be used to adjust the flexibility of the trend and to choose how aggressively the model should follow historical trend changes. Analysts may also adjust the number of potential changepoints included or manually

specify their location. This allows non-experts with knowledge about events that may affect the growth rate to use the parameter as a knob to either increase or decrease the number of changepoints included (Taylor and Letham 2017). It also allows for the analyst to add changepoints which the automatic selection procedure may have missed or remove changepoints when the model is overfitting historical trends (Taylor and Letham 2017).

The decomposable form of the model allows for multiple seasonality components with different periods to be added to the model. Seasonality components are modelled by a Fourier series and have a Normal prior distribution placed on its parameters (Taylor and Letham 2017). The spread parameter can be adjusted by analysts to smooth the seasonality and change how much of historical seasonality is projected into the future (Taylor and Letham 2017).

The analyst may also provide a list of important events and holidays which have impacted the time series in the past or which they know might impact it in the future. The list could include the name, date, and country in which they have taken place or are expected to take place (Taylor and Letham 2017). By specifying the country of occurrence, separate lists can be populated for global events and holidays, and country-specific events and holidays. The union of the global and country-specific lists can then be used for forecasting. Like seasonality, a Normal prior distribution is placed on the parameters of the holiday component, and the scale parameter can be adjusted by analysts to smooth the holidays (Taylor and Letham 2017).

Prophet’s Bayesian approach to forecasting allows the analyst to incorporate their expert knowledge into the model building process and has produced significantly improved forecasts compared to the ARIMA model. Taylor and Letham (2017) forecasted the number of events on Facebook using Prophet.

The time series was impacted by holidays, had strong multi-period seasonality, and a piecewise trend. The forecasts produced by Prophet were compared to common forecasting techniques such as exponential smoothing, ARIMA, the seasonal nave, and the nave model. While exponential smoothing and the seasonal nave model were quite robust, the ARIMA forecasts were fragile. No model besides Prophet accounted for the dips around holidays and the upward trend of the time series towards later observations.

Prophet’s ability to forecast at scale enables it to model a wide variety of data and may be able to adequately fit exchange rate data. Its non-linear components could capture the non-linearity present in exchange rates. In contrast, the ARIMA model is a linear function of previous observations and lagged residuals and fails to capture the non-linearity inherent in exchange rate data. Prophet performs well on data with strong multiple “human scale” seasonalities and historical trend changes. This differs to the ARIMA model which requires the data to be de-trended and the variance stabilised before the model can be fitted. Hence Prophet could model the weekly seasonality of closing prices due to low trading activity which occurs around the weekend and high trading activity which occurs mid-week.

Choosing the correct combination of parameters for the ARIMA model is a challenging task due to the array of possible choices. Although the `auto.arima` function in R may be used to automatically select an ARIMA model that best fits the data, completely automatic forecasting methods are too brittle and do not allow for useful assumptions to be incorporated into the model. Prophet makes use of a semi-automatic forecasting technique that keeps the analyst-in-the-loop. Its default settings are said to generate forecasts that are as accurate as those produced by skilled forecasters. If the forecasts produced are unsatisfactory, they can be improved by the analyst

by configuring the model through its easily interpretable parameters. Furthermore, non-experts who have domain knowledge about factors that affect Bitcoin or if the dates of events which could impact the price of Bitcoin are known, it may be incorporated into the model by the analyst. Prophet can produce forecasts over irregular time intervals and allows for missing values in the time series without the need for interpolation (Taylor and Letham 2017). The ARIMA model on the other hand requires large outliers to be removed and handles missing values by interpolation. If Prophet produces forecasts that are as good as the ARIMA model when forecasting BTC/ZAR, it can be compared to more complex forecasting methods that are less tractable such as the hybrid models and machine learning techniques seen earlier.

3 Data

The data set used in this study consists of the daily closing prices of BTC/ZAR over the period 24 January 2016 to 17 July 2017. This comprises of a total of 541 trading days and was the chosen time period for analysis due to constraints in obtaining data over a longer period. The data was obtained from Bitcoincharts.

Bitcoin is of specific interest as it presents an interesting parallel to traditional exchange rate markets. The cryptocurrency is built on a decentralised system and as a result its value cannot be directly influenced by a central authority (Fantazzini and Nigmatullin 2016). In addition, the Bitcoin market has attracted attention worldwide and is currently the leading cryptocurrency, with awareness and adoption of the currency growing over time (Fantazzini and Nigmatullin 2016). Its novelty makes it a highly volatile hence speculative market and provides an opportunity for forecasting.

In this study, the analysis of BTC/ZAR forecasts are based on the daily continuous log returns. The daily closing prices are transformed into returns by taking the log difference at each time t and is calculated as follows:

$$r_t = \ln\left(\frac{p_t}{p_{t-1}}\right)$$

where p_t represents the closing price and r_t represents the log return for time $t = 1, 2, \dots, T$.

4 Methodology

In this paper, the out-of-sample forecasts of ARIMA and Prophet are compared over varying forecast horizons. The data is first split into a training set and test set. The training set starts on the 24th January 2016 and ends on the 16th January 2017, while the test set starts on the 17th January 2017 and ends on the 17th July 2017. The Box-Jenkins methodology is then used to select the correct ARIMA model while Prophet’s semi-automatic procedure is used for model selection. The selected ARIMA and Prophet model are then fitted to the training set and rolling window forecasts are made 1 day, 30 days and 90 days ahead. This allows us to examine the forecast horizon effect. The Mincer-Zarnowitz test is then used to evaluate the forecasts produced by ARIMA and Prophet against the test set. Since there is no consensus on which accuracy statistic best measures the performance of forecasting techniques, the most popular criteria, the Mean Squared Error (MSE) is employed. The results obtained from the ARIMA model will be used as a benchmark for comparison. This process is then repeated based on data which has outliers and missing values present.

4.1 The Box-Jenkins Methodology

Box and Jenkins proposed a set of guidelines that can be followed when selecting ARIMA models. This systematic procedure to designing ARIMA models has made them highly popular (Hibon and Makridakis 1997). It consists of a four-stage iterative process in which:

1. The process is either transformed or differenced to de-trend and stabilise the variance of the data.
2. The autocorrelation and partial autocorrelation plots are used to determine the order of p and q .
3. The parameters of the model are then estimated.
4. A diagnostic check is performed to ensure that the residuals are a white noise process.

If the residuals are not white noise, steps 2-4 are repeated until a satisfactory model is identified. On the contrary, if the diagnostic check shows that the residuals are random, the developed model will be the final model used for forecasting.

4.2 Semi-Automatic Selection

When a large number of forecasts are produced, manually identifying problematic forecasts becomes a time consuming and difficult task. Prophet provides a semi-automated forecast evaluation system that selects the best model which fits the data. When there are large forecast errors, the forecasts are flagged so that the analyst can explore the cause of the errors, identify and remove potential outliers and either adjust the model or choose a more

appropriate model (Taylor and Letham 2017). This keeps the analyst-in-the-loop. Prophet has the following default settings:

- The trend is set to be linear
- The width of the uncertainty intervals is set to 80%
- Weekly and yearly seasonality are automatically detected and included in the model if present
- The smoothing parameter for holidays and seasonality is set at 10
- The smoothing parameter for the trend is set to 0.05
- The number of potential changepoints is set to 25

This paper makes use of Prophet’s default settings to fit the data, however weekly seasonality is included in the model. This allows us to account for our prior knowledge about the closing prices of Bitcoin which tend to be lower around the weekend however higher mid-week, due to fluctuations in trading activity. Furthermore, a linear trend is appropriate since Bitcoin does not have an upper limit on its closing prices.

4.3 Rolling Window Forecasts

Rolling window forecasts are useful in evaluating the robustness of a forecasting method. In a rolling window forecast, the forecasts are made h -steps at a time and the actual observation rather than the predicted value is used for the next prediction in the forecast horizon. In this way, a poor forecast will not have negative consequences on future forecasts to be made, since the observed values will be used to correct itself for the remaining forecasts (Zhang and Hu 1998). In this paper, we first fit the ARIMA and Prophet

model to the training set and forecast 1 day ahead. We then refit the model to the training set, increasing its size by one observation, and the next 1 day ahead forecast is made. This process is repeated until all forecasts have been made into the test set. The 1 day ahead forecasts are then compared to the 1 day ahead observed values in the test set and the error statistics are calculated. This procedure is then repeated for the 30 day and 90 day ahead forecasts.

4.4 Mincer-Zarnowitz Approach to Forecast Evaluation

The Mincer-Zarnowitz approach to evaluating forecast accuracy is commonly used and can be useful when comparing the forecasts produced by Prophet and the ARIMA model. Mincer and Zarnowitz (1969) proposed an absolute measure to evaluate forecast accuracy which considers the distance between the actual and predicted values. To analyse the absolute errors produced by the forecasts, the observed values r_t are regressed against the predicted values \hat{r}_t , i.e.

$$r_t = \alpha + \beta \cdot \hat{r}_t + e_t$$

where α represents the mean distance between the observed and predicted values while β represents the correlation between the forecasted errors and the predicted values. When $\alpha = 0$ it implies that the forecasts are unbiased and do not systematically overestimate or underestimate the data. When $\beta = 1$ it implies that the forecasts are efficient and uncorrelated with the forecasted errors. A joint hypothesis test of $H_0 : \alpha = 0 \cup \beta = 1$ is performed to check the efficiency and bias of the forecasts and the model which produces the best results generates superior forecasts.

4.5 Evaluation of forecasts based on data with missing values and outliers

The original dataset is modified in order to examine how robust Prophet and ARIMA are when fitted to data with missing values and outliers. 20 values are randomly removed from the training set and 5 outliers are randomly inserted in place of existing values. The models are then fitted to this “dirty” dataset and evaluating of forecasts are carried out as before.

5 Results

5.1 Evaluation over varying forecast horizons

Although the Box-Jenkins methodology resulted in an ARIMA(0,0,1) model being the best fit to the training set, an ARIMA(1,0,0) is employed as the benchmark for comparison so that a rolling window forecast may be conducted. The AIC of the ARIMA(1,0,0) is marginally higher than that of the ARIMA(0,0,1) and analysis of the p-value indicates that the AR parameter is statistically significant in the explanation of the movement of BTC/ZAR. The Prophet model fitted to the training set included a linear trend and weekly seasonality, with all other parameters set to the default values as described in section 4.2. The automatic-selection procedure for ARIMA and Prophet selected the same model as those chosen manually.

5.1.1 1-step ahead Forecasts

Figure 1 depicts the resulting forecasts from fitting both models. The forecasts produced by the ARIMA model and Prophet lie within a band centred around zero. Both models fail to capture the large movement of returns

away from zero. The 80% uncertainty intervals produced by both models are approximately the same width and begin to widen as forecasts are made further into the test set.




Figure 5.1: 1-step ahead forecasts from the ARIMA(1,0,0) and Prophet model

Table 1 displays the MSE and the Mincer-Zarnowitz results obtained from regressing the test set returns against the one-step ahead forecasts. The intercept estimate is small and positive for ARIMA and Prophet, with the intercept for ARIMA being closer to zero than Prophet. This indicates that both models systematically underestimate the 1-step ahead returns, with ARIMA yielding mean forecasts closer to the observed mean. The slope estimate is negative for both models indicating that the observed returns are actually the opposite of the suggested forecast. The F-test for the joint

hypothesis of a unity slope and zero intercept is rejected at the 5% significance level for both models, with the p-value for ARIMA being much larger than Prophet. Hence the forecasts generated by both ARIMA and Prophet are biased and/or inefficient. The MSE produced by both models agree with the Mincer-Zarnowitz results, with ARIMA producing a marginally smaller MSE than Prophet.

	intercept ↕	slope ↕	p-value ↕	MSE ↕
ARIMA	0.009545697	-2.005624	0.025448831	0.001941537
Prophet	0.013701904	-1.709744	0.007152494	0.001958101

Figure 5.2: Mincer-Zarnowitz results for 1-step ahead forecasts

5.1.2 30-steps ahead Forecasts

Figure 2 illustrates the resulting 30-step ahead forecasts produced from fitting ARIMA and Prophet. As in the 1-step ahead case, both models fail to capture the large movement of returns away from zero with its forecasts lying within a band centred around zero. The 80% uncertainty intervals produced by both models are approximately the same width and marginally wider than the uncertainty intervals produced by the 1-step ahead forecasts.




Figure 5.3: 30-step ahead forecasts from the ARIMA(1,0,0) and Prophet model

Table 2 displays the MSE and the Mincer-Zarnowitz results obtained from regressing the test set returns against the 30-step ahead forecasts. The intercept estimate is small and positive for ARIMA and Prophet, with a smaller intercept estimated for Prophet. This indicates that both models systematically underestimate the 30-step ahead returns, with Prophet yielding mean forecasts closer to the observed mean. The slope estimate is large and negative for both models with the slope estimated to be as low as -12.74 for ARIMA. This suggests that the observed returns are actually the opposite of the proposed forecast. The F-test for the joint hypothesis of a unity slope and zero intercept is rejected at the 5% significance level for Prophet only. The p-value of 0.0614 from the joint hypothesis test is slightly significant for

ARIMA. Hence the forecasts generated by Prophet are biased and/or inefficient while the forecasts generated by ARIMA are unbiased and/or efficient. The MSE produced by both models agree with the Mincer-Zarnowitz results, with ARIMA producing a marginally smaller MSE than Prophet.

	intercept $\hat{\diamond}$	slope $\hat{\diamond}$	p value $\hat{\diamond}$	MSE $\hat{\diamond}$
ARIMA	0.03132042	-12.736757	0.0613586071	0.001988463
Prophet	0.01716876	-2.621914	0.0003772506	0.002075491

Figure 5.4: Mincer-Zarnowitz results for 30-step ahead forecasts

5.1.3 90-steps ahead Forecasts

Figure 3 illustrates the 90-step ahead forecasts produced by ARIMA and Prophet. Both models fail to capture large movement of returns away from zero as in the case of the 1-step ahead and 30-step ahead forecasts. The 80% uncertainty intervals produced by both models are approximately the same width and are marginally narrower than the uncertainty intervals produced by the one-step ahead and 30-step ahead forecasts.

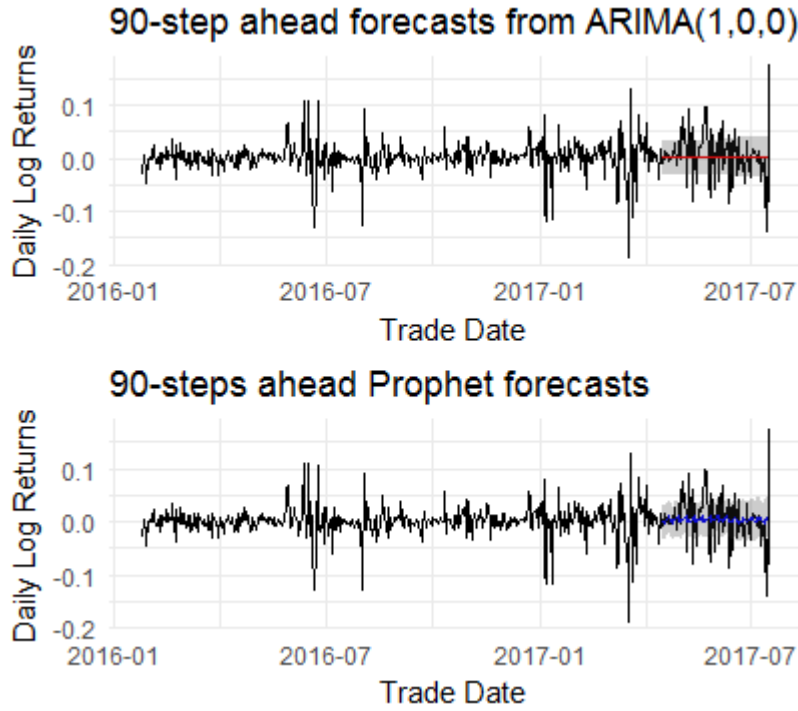


Figure 5.5: 90-step ahead forecasts from the ARIMA(1,0,0) and Prophet model

Table 3 presents the MSE and the Mincer-Zarnowitz results obtained from regressing the test set returns against the 90-step ahead forecasts. The intercept estimate is very small and positive for ARIMA and Prophet, with a marginally smaller intercept estimated for ARIMA. This indicates that both models systematically underestimate the 90-step ahead returns, with ARIMA yielding mean forecasts closer to the observed mean. The slope estimate is small and negative for both models with the slope estimated for ARIMA lying closer to -1. This suggests that the observed 90-steps ahead returns are actually the opposite of the proposed 90-steps ahead forecasts. The F-test for the joint hypothesis of a unity slope and zero intercept fails to be rejected at the 5% significance level for both ARIMA and Prophet,

with the p-value from the test being much larger for ARIMA. Hence the forecasts generated by Prophet and ARIMA are unbiased and/or efficient. The MSE produced by both models agree with the Mincer-Zarnowitz results, with ARIMA producing a marginally smaller MSE than Prophet.

	intercept ↕	slope ↕	p value ↕	MSE ↕
ARIMA	0.006012467	-0.9386872	0.8430658	0.001790129
Prophet	0.006953398	-0.6154348	0.4086752	0.001813925

Figure 5.6: Mincer-Zarnowitz results for 90-step ahead forecasts

5.1.4 Comparison over varying forecast horizons

Table 4 and 5 compares the p-values obtained from the Mincer-Zarnowitz tests and the MSE of the forecasts produced by ARIMA and Prophet over the three forecast horizons. In comparing the forecasting accuracy of the ARIMA model over the three forecast horizons, we observe that the p-value gets more significant as the forecast horizon gets longer. This is in contrast to the MSE of the ARIMA forecasts which is smaller at the shortest and longest forecast horizons. In comparing the forecasting accuracy of Prophet, we find that the p-value is more significant for the shortest and longest forecast horizon. This agrees with the MSE of the Prophet forecasts which is smaller at the longest and shortest forecast horizon. Both ARIMA and Prophet generate forecasts with the smallest MSE and the most significant p-value at the longest forecast horizon. In comparing the forecasting accuracy of ARIMA and Prophet over the three forecast horizons, we notice that the p-value obtained from the Mincer-Zarnowitz results is more significant for ARIMA

than Prophet at every forecast horizon. This echoes the results obtained from the MSE which is marginally smaller for ARIMA than Prophet over all forecast horizons.

	one_step $\hat{\Delta}$	thirty_steps $\hat{\Delta}$	ninety_steps $\hat{\Delta}$
ARIMA	0.025448831	0.0613586071	0.8430658
Prophet	0.007152494	0.0003772506	0.4086752

Figure 5.7: Mincer-Zarnowitz results for all forecast horizons

	one_step $\hat{\Delta}$	thirty_steps $\hat{\Delta}$	ninety_steps $\hat{\Delta}$
ARIMA	0.001941537	0.001988463	0.001790129
Prophet	0.001958101	0.002075491	0.001813925

Figure 5.8: MSE for all forecast horizons

5.2 Evaluation on data with missing values and outliers present

By following the Box-Jenkins methodology for selection of an ARIMA model, an ARIMA(3,0,0) was employed to fit the training set which included missing values and outliers. Although the ARIMA(0,0,3) had a slightly smaller AIC, a decision was made to use the ARIMA(3,0,0) instead so that a rolling window forecast may be used. Analysis of the AR parameters indicate that they are statistically significant in the explanation of the movement of BTC/ZAR. The Prophet model fitted to the training set included a linear trend and weekly seasonality. All other parameters were set to Prophet's default values

as described in section 4.2. The automatic-selection procedure for Prophet and ARIMA selected the same models as those chosen manually.

5.2.1 1-step ahead Forecasts

Table 6 shows the MSE of the forecasts generated by ARIMA and Prophet and the Mincer-Zarnowitz results obtained from regressing the test set returns against the 1-step ahead returns forecasted. We observe that a positive intercept is estimated in the case of ARIMA and Prophet, indicating that both models systematically underestimate the 1-step ahead returns. The intercept estimated for ARIMA is smaller than Prophet, indicating that the mean forecasted returns yielded by ARIMA is closer to the mean observed returns. The intercept estimates of 0.011895 and 0.014802 for ARIMA and Prophet respectively are larger than the intercept estimates of 0.009546 and 0.013702, which was based on the “clean” dataset. This suggests that both models systematically underestimates the returns more when there are missing values and outliers present in the data.

We also observe that the slope estimate is small and negative for both models with the slope estimated for ARIMA lying closer to -1. This suggests that the observed 1-step ahead returns are actually the opposite of the proposed 1-step ahead forecasts. The slope estimates of -0.883228 and -0.676132 for ARIMA and Prophet respectively are less steep than the slope estimates of -2.005624 and -1.709744, which was based on the evaluation of data that had no values missing or outliers present.

	intercept ↕	slope ↕	p-value ↕	MSE ↕
ARIMA	0.01189529	-0.8832276	3.182484e-05	0.002072583
Prophet	0.01480189	-0.6761324	5.315724e-06	0.002135501

Figure 5.9: Mincer-Zarnowitz results for 1-step ahead forecasts

The p-values obtained from the F-test for the joint hypothesis of a unity slope and zero intercept is approximately zero for ARIMA and Prophet, with the p-value for Prophet lying closer to zero than ARIMA. The null hypothesis is thus rejected at the 1% significance level, indicating that the forecasts generated by Prophet and ARIMA are biased and/or inefficient. The MSE produced by both models agree with the Mincer-Zarnowitz results, with ARIMA producing a marginally smaller MSE than Prophet. The results obtained from evaluating the forecasting accuracy of ARIMA and Prophet based on the “clean” dataset show similar findings, however led to a smaller MSE of 0.001942 and 0.001958 and a more significant p-value of 0.025449 and 0.007152 for ARIMA and Prophet respectively.

5.2.2 30-steps ahead Forecasts

Table 7 shows the MSE of the forecasts produced by ARIMA and Prophet and the Mincer-Zarnowitz results obtained from regressing the test set returns against the 30-step ahead returns forecasted. We observe that a positive intercept is estimated in the case of ARIMA and Prophet, indicating that both models systematically underestimate the 30-step ahead returns. The intercept estimated for ARIMA is larger than Prophet, indicating that the mean forecasted returns yielded by Prophet is closer to the mean observed

returns. The intercept estimate of 0.11414 for ARIMA is larger than the estimate of 0.031320 which was based on the evaluation of the “clean” dataset. This suggests that the ARIMA model underestimates the returns more when there are missing values and outliers present in the data. The intercept estimate of 0.014396 for Prophet is less than the estimate of 0.017169 which was based on the evaluation of data that had no values missing or outliers present. This suggests that Prophet underestimates the returns less when there are missing values and outliers present in the data.

	intercept $\hat{\diamond}$	slope $\hat{\diamond}$	p-value $\hat{\diamond}$	MSE $\hat{\diamond}$
ARIMA	0.11414036	-12.9853346	1.075322e-01	0.002003289
Prophet	0.01439573	-0.6516095	1.479928e-06	0.002321602

Figure 5.10: Mincer-Zarnowitz results for 30-step ahead forecasts

We also observe that the slope estimate is large and negative for ARIMA and small and negative for Prophet. This suggests that the observed 30-step ahead returns are actually the opposite of the proposed 30-step ahead forecasts. The slope estimate of -12.985335 for ARIMA is more steep than the slope estimate of -12.736757 which was based on the evaluation of data that had no values missing or outliers present. This suggests that the efficiency of the forecasted returns deteriorates when there are missing values and outliers present in the data. The slope estimate of -0.65161 for Prophet is less steep than the estimate -2.621914, which was based on the evaluation of data that had no values missing or outliers present. This suggests that the efficiency of the forecasted returns improve when there are missing values and outliers present in the data.

The p-value obtained from the F-test for the joint hypothesis of a unity slope and zero intercept is approximately zero for Prophet. The null hypothesis is thus rejected at the 1% significance level, indicating that the 30-step ahead forecasts generated by Prophet are biased and/or inefficient. The p-value of 0.107532 for ARIMA is significant at the 5% level, suggesting at the 30-step ahead forecasts produced by ARIMA are unbiased and/or efficient. MSE produced by both models agree with the Mincer-Zarnowitz results, with ARIMA producing a smaller MSE than Prophet. The results obtained from evaluating the forecasting accuracy of ARIMA and Prophet based on the “clean” dataset, led to a more significant p-value of 0.000377 for Prophet only, and a smaller MSE of 0.001988 and 0.002075 for ARIMA and Prophet respectively.

5.2.3 90-steps ahead Forecasts

Table 8 shows the MSE of the forecasts produced by ARIMA and Prophet and the Mincer-Zarnowitz results obtained from regressing the test set returns against the 90-step ahead returns forecasted. We observe that a negative intercept is estimated in the case of ARIMA and Prophet, indicating that both models systematically overestimate the 90-step ahead returns. The intercept estimated for Prophet lies closer to zero, indicating that the mean forecasted returns yielded by Prophet is closer to the mean observed returns. This is in contrast to the positive intercepts that were estimated based on the data that had no values missing or outliers present.

We also observe that the slope estimate is large and positive for ARIMA however small and positive for Prophet. This is in contrast to the small and negative slope estimates which were based on the evaluation of data that had no values missing or outliers present.

	intercept $\hat{\Delta}$	slope $\hat{\Delta}$	p-value $\hat{\Delta}$	MSE $\hat{\Delta}$
ARIMA	-0.226672873	27.0635249	0.0129409634	0.001789444
Prophet	-0.009290786	0.6438489	0.0007635404	0.002030498

Figure 5.11: Mincer-Zarnowitz results for 90-step ahead forecasts

The p-values obtained from the F-test for the joint hypothesis of a unity slope and zero intercept are insignificant for ARIMA and Prophet, with the p-value for Prophet lying close to zero. The null hypothesis is thus rejected at the 5% significance level, indicating that the 90-steps ahead forecasts generated by Prophet and ARIMA are biased and/or inefficient. The MSE produced by both models agree with the Mincer-Zarnowitz results, with ARIMA producing a marginally smaller MSE than Prophet. The results obtained from evaluating the forecasting accuracy of ARIMA and Prophet based on data which had no missing values or outliers present contrasts these findings, with a p-value of 0.843066 and 0.408675 for ARIMA and Prophet respectively.

5.2.4 Comparison over varying forecast horizons

Table 9 and 10 compares the p-values obtained from the Mincer-Zarnowitz tests and the MSE of the forecasts produced by ARIMA and Prophet over the three forecast horizons. In comparing the forecasting accuracy of the ARIMA model over the three forecast horizons, we observe that the p-value is more significant at the longer forecast horizons. This agrees with the MSE of the ARIMA forecasts which is smaller at the longer forecast horizons. In comparing the forecasting accuracy of Prophet, we find that the p-value is

more significant for the longest forecast horizon. This agrees with the MSE of the Prophet forecasts which is smaller at the longest forecast horizon.

	<u>One step</u>	<u>Thirty steps</u>	<u>Ninety steps</u>
ARIMA	3.182484e-05	1.075322e-01	0.0129409634
Prophet	5.315724e-06	1.479928e-06	0.0007635404

Figure 5.12: Mincer-Zarnowitz results for all forecast horizons

	<u>One step</u>	<u>Thirty steps</u>	<u>Ninety steps</u>
ARIMA	0.002072583	0.002003289	0.001789444
Prophet	0.002135501	0.002321602	0.002030498

Figure 5.13: MSE for all forecast horizons

In comparing the forecasting accuracy of ARIMA and Prophet over the three forecast horizons, we notice that the p-value obtained from the Mincer-Zarnowitz results is more significant for ARIMA than Prophet at every forecast horizon. This echoes the results obtained from the MSE which is marginally smaller for ARIMA than Prophet over all forecast horizons.

6 Discussion and Conclusions

The results show a clear difference in the effect of the forecast horizon in the case of the clean data. The results based on the MSE and the p-value obtained from the Mincer-Zarnowitz test surprisingly show that the ARIMA model produces superior forecasts of the returns at the longest forecast horizon. Similarly, Prophet produces superior forecasts of the returns at

the longest forecast horizon. We notice that although the ARIMA model outperforms Prophet at every forecast horizon, the difference between the MSE's of the forecasts generated by both models is insignificant. Furthermore, the Mincer-Zarnowitz p-values for both models are significant at the same forecast horizons, suggesting that the efficiency and biasness of the return forecasts yielded by ARIMA and Prophet are synchronized at each forecast horizon. If a small amount of predictive accuracy is willing to be sacrificed by the analyst in the forecasting of BTC/ZAR, it might be useful to employ Prophet instead of ARIMA. This would enable traders with expert knowledge about Bitcoin, however no statistical knowledge about forecasting models, to incorporate their domain knowledge into the Prophet model through its intuitively adjustable parameters. Furthermore, as more information about events that may affect the price of Bitcoin becomes available, it may be included in the model and could potentially change the predictive accuracy of Prophet.

The results based on the unclean data show that both ARIMA and Prophet are sensitive to the presence of outliers and missing values in addition to the forecast horizon. We observe that Prophet yields superior forecasts at the longest forecast horizon. Although the MSE suggests that the ARIMA model produces forecasts with minimal error at the longest forecast horizon, analysis of the p-value indicates that the forecasts are efficient and/or unbiased at 30-steps ahead only. It is evident that Prophet produces inferior results in comparison to the ARIMA model over all forecast horizons. This is unexpected as Prophet is said to be robust to the presence of outliers and missing values in the data, without the need for interpolation. We also notice that the forecasts produced by both models deteriorates when there are missing values and outliers present in the data, with the only efficient and/or

unbiased forecasts yielded by the ARIMA model at a forecast horizon of 30 days.

The limitation of this study is that it looks specifically at BTC/ZAR. Different results may emerge for fiat currencies or the exchange rate between Bitcoin with other currencies. Hence caution should be exercised when generalising these results to other exchange rates. As a point of further research, more complex techniques such as neural networks could be used as a benchmark for comparison when evaluating the forecasts produced by Prophet.

References

- Abu-Mostafa, Yaser S, and Amir F Atiya. 1996. "Introduction to Financial Forecasting." *Applied Intelligence* 6 (3). Springer: 205–13.
- Anastasakis, Leonidas, and Neil Mort. 2009. "Exchange Rate Forecasting Using a Combined Parametric and Nonparametric Self-Organising Modelling Approach." *Expert Systems with Applications* 36 (10). Elsevier: 12001–11.
- Castillo, Oscar, and Patricia Melin. 2002. "Hybrid Intelligent Systems for Time Series Prediction Using Neural Networks, Fuzzy Logic, and Fractal Theory." *IEEE Transactions on Neural Networks* 13 (6). IEEE: 1395–1408.
- Cheung, Yin-Wong. 1993. "Long Memory in Foreign-Exchange Rates." *Journal of Business & Economic Statistics* 11 (1). Taylor & Francis: 93–101.
- Etuk, Ette Harrison, Dagogo SA Wokoma, and Imoh Udo Moffat. 2013. "Additive Sarima Modelling of Monthly Nigerian Naira-Cfa Franc Exchange Rates." *European Journal of Statistics and Probability* 1 (1): 1–12.
- Fahimifard, SM, Masuod Homayounifar, M Sabouhi, and AR Moghadamnia. 2009. "Comparison of Anfis, Ann, Garch and Arima Techniques to Exchange Rate Forecasting." *Journal of Applied Sciences* 9 (20): 3641–51.
- Hibon, Michael, and Spyros Makridakis. 1997. "ARMA Models and the Box–Jenkins Methodology." John Wiley & Sons, Ltd.
- Ince, Huseyin, and Theodore B Trafalis. 2006. "A Hybrid Model for Exchange Rate Prediction." *Decision Support Systems* 42 (2). Elsevier: 1054–62.
- Kamruzzaman, Joarder, and Ruhul A Sarker. 2003. "Forecasting of Currency Exchange Rates Using Ann: A Case Study." In *Neural Networks and Signal Processing, 2003. Proceedings of the 2003 International Conference on*, 1:793–97. IEEE.
- Khashei, Mehdi, and Hehdi Bijari. 2011. "Exchange Rate Forecasting

Better with Hybrid Artificial Neural Networks Models.” *Journal of Mathematical and Computational Science* 1 (1). Science & Knowledge Publishing Corporation Limited (SCIK): 103.

Khashei, Mehdi, Mehdi Bijari, and Gholam Ali Raissi Ardali. 2009. “Improvement of Auto-Regressive Integrated Moving Average Models Using Fuzzy Logic and Artificial Neural Networks (Anns).” *Neurocomputing* 72 (4). Elsevier: 956–67.

———. 2012. “Hybridization of Autoregressive Integrated Moving Average (Arima) with Probabilistic Neural Networks (Pnns).” *Computers & Industrial Engineering* 63 (1). Elsevier: 37–45.

Lin, Chiun-Sin, Sheng-Hsiung Chiu, and Tzu-Yu Lin. 2012. “Empirical Mode Decomposition-based Least Squares Support Vector Regression for Foreign Exchange Rate Forecasting.” *Economic Modelling* 29 (6). Elsevier: 2583–90.

Mincer, Jacob A, and Victor Zarnowitz. 1969. “The Evaluation of Economic Forecasts.” In *Economic Forecasts and Expectations: Analysis of Forecasting Behavior and Performance*, 3–46. NBER.

Nwankwo, Steve C. 2014. “Autoregressive Integrated Moving Average (Arima) Model for Exchange Rate (Naira to Dollar).” *Academic Journal of Interdisciplinary Studies* 3 (4): 429.

Santos, Andr Alves Portela, Newton Carneiro Affonso da Costa, and Leandro dos Santos Coelho. 2007. “Computational Intelligence Approaches and Linear Models in Case Studies of Forecasting Exchange Rates.” *Expert Systems with Applications* 33 (4). Elsevier: 816–23.

Taylor, Sean J, and Benjamin Letham. 2017. “Forecasting at Scale.”

Tseng, Fang-Mei, Gwo-Hshiung Tzeng, Hsiao-Cheng Yu, and Benjamin JC Yuan. 2001. “Fuzzy Arima Model for Forecasting the Foreign Exchange

Market.” *Fuzzy Sets and Systems* 118 (1). Elsevier: 9–19.

Zhang, Gioqinang, and Michael Y Hu. 1998. “Neural Network Forecasting of the British Pound/Us Dollar Exchange Rate.” *Omega* 26 (4). Elsevier: 495–506.